

Practical Machine Learning - Week4

Saul Lugo

January 31, 2016

Reguralized Regression

Model Selection Approach: split samples

1. Divide data into training/test/validation sets
2. Treat validation as a test data, train all competing models on the train data and pick the best one on validation.
3. To appropriately assess performance on new data apply the model to test set
4. You might re-split and reperform steps 1-3

Combining Predictors

The idea is to combine different predictor models in order to improve the accuracy. The following example combines a GLM with a Random Forest predictor:

```
library(ISLR); data(Wage); library(ggplot2); library(caret);
Wage <- subset(Wage, select=-c(logwage))

set.seed(1234)
#Splitting the data into training, testing and validation sets
inBuild <- createDataPartition(y=Wage$wage,p=0.7,list=FALSE)
validation <- Wage[-inBuild,]; buildData <- Wage[inBuild,]
inTrain <- createDataPartition(y=buildData$wage,p=0.7,list=FALSE)
training <- buildData[inTrain,]; testing <- buildData[-inTrain,]

dim(training); dim(testing); dim(validation);

## [1] 1474    11

## [1] 628    11

## [1] 898    11

#Now, fit the two models
mod1 <- train(wage ~.,method="glm",data=training)
mod2 <- train(wage ~.,method="rf",data=training,trainControl=trainControl(method="cv",number=3))

mod1

## Generalized Linear Model
##
## 1474 samples
##    10 predictor
```


```
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 1474, 1474, 1474, 1474, 1474, 1474, ...
## Resampling results
##
##      RMSE      Rsquared    RMSE SD    Rsquared SD
##      35.0092   0.3204793   1.762416   0.02920338
##
##
```

```
mod2
```

```
## Random Forest
##
## 1474 samples
##   10 predictor
##
## No pre-processing
## Resampling: Cross-Validated (3 fold)
## Summary of sample sizes: 982, 983, 983
## Resampling results across tuning parameters:
##
##      mtry  RMSE      Rsquared    RMSE SD    Rsquared SD
##      2     38.16546  0.3123716  3.278744  0.02009053
##      13     37.05130  0.2644913  1.636132  0.03221683
##      25     38.47024  0.2321406  1.272379  0.03441396
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 13.
```

```
#plotting mod1 vs mod2
```

```
pred1 <- predict(mod1,testing); pred2 <- predict(mod2,testing);
qplot(pred1,pred2,colour=wage,data=testing)
```



```
#Fit a model that combines both predictors
predDF <- data.frame(pred1,pred2,wage=testing$wage)
combModFit <- train(wage ~.,method="gam",data=predDF)
combPred <- predict(combModFit,predDF)
combModFit
```

```
## Generalized Additive Model using Splines
##
## 628 samples
## 2 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 628, 628, 628, 628, 628, 628, ...
## Resampling results across tuning parameters:
##
## select RMSE Rsquared RMSE SD Rsquared SD
## FALSE 33.64839 0.3406428 2.142237 0.04548650
## TRUE 33.68495 0.3392327 2.145042 0.04566066
##
## Tuning parameter 'method' was held constant at a value of GCV.Cp
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were select = FALSE and method
## = GCV.Cp.
```

```
#Comparing the Root Squared Errors between mod1, mod2 and the combined model
sqrt(sum(pred1-testing$wage)^2)
```

```
## [1] 566.8475
```

```
sqrt(sum(pred2-testing$wage)^2)
```

```
## [1] 528.6399
```

```
sqrt(sum(combPred-testing$wage)^2)
```

```
## [1] 1.336176e-10
```

```
#Checking the model on the validation set  
pred1V <- predict(mod1,validation); pred2V <- predict(mod2,validation)  
predVDF <- data.frame(pred1=pred1V,pred2=pred2V)  
combPredV <- predict(combModFit,predVDF)
```

```
#Checking the RSE on the validation set  
sqrt(sum(pred1V-validation$wage)^2)
```

```
## [1] 342.1366
```

```
sqrt(sum(pred2V-validation$wage)^2)
```

```
## [1] 285.9747
```

```
sqrt(sum(combPredV-validation$wage)^2)
```

```
## [1] 656.5295
```

Unsupervised Prediction

When performing unsupervised prediction one does not know the labels of the outcome beforehand.

The following example use clustering technique on the IRIS dataset. We ignore the Species variable in order to simulate that we don't know the outcome. We use clustering by **k-means** in order to build the clusters, then we build a prediction model using the clusters in the training dataset as the outcome variable:


```
data(iris)  
library(ggplot2)  
library(caret)  
  
inTrain <- createDataPartition(y=iris$Species,p=.7,list=FALSE)  
training <- iris[inTrain,]  
testing <- iris[-inTrain,]  
dim(training); dim(testing)
```

```
## [1] 105 5
```

```
## [1] 45 5
```

```
#Building the clusters
```

```
kMeans1 <- kmeans(subset(training,select=-c(Species)),centers=3)
training$clusters <- as.factor(kMeans1$cluster)
qplot(Petal.Width,Petal.Length,colour=clusters,data=training)
```



```
table(kMeans1$cluster,training$Species)
```

```
##
##      setosa versicolor virginica
##  1         0          34         8
##  2         0           1        27
##  3        35           0         0
```

```
#Fitting the model using the clusters
```

```
modFit <- train(clusters ~.,data=subset(training,select=-c(Species)),method="rpart")
table(predict(modFit,training),training$Species)
```

```
##
##      setosa versicolor virginica
##  1         0          35         8
##  2         0           0        27
##  3        35           0         0
```

```
#Apply the model on the test set
```

```
testClusterPred <- predict(modFit,testing)
table(testClusterPred,testing$Species)
```

```
##
## testClusterPred setosa versicolor virginica
##          1      0      15      8
##          2      0       0      7
##          3     15       0      0
```